

# User manual

Compiling and cleaning

From the top level project directory:

- `make` - builds all parts (P1, P2, P3, P4\_P5)
- `make clean` - removes the compiled executables

P1 (cd p1)

- `./reverse_server`
- `./reverse_client <server_ip> <text...>`

P2 (cd p2)

- `./ls_server`
- `./ls_client <server_ip> <ls_args...>`

P3 (cd p3)

- `./disk_server <cylinders> <sectors> <track_delay_us> <backing_file>`
- `./disk_client <server_ip>`
- `./disk_rand <server_ip> <ops> <seed>`

P4\_P5 (cd p4\_p5)

- `./file_system_server <disk_server_ip>`
- `./file_system_client <fs_server_ip>`

# Technical manual

## Disk design (Part 3)

The disk server simulates a block device using one normal file as the “disk”.

Data layout:

- Cylinders (C) and sectors per cylinder (S) are passed args from the command line.
- Block size is 128 bytes. Total blocks =  $C \times S$ .
- Each block is addressed by (cylinder c, sector s).
- Block number  $b = c \times S + s$ , and byte offset in the backing file =  $b \times 128$ .
- Every disk read/write is exactly 128 bytes at that offset.

In-memory state per client:

- `current_cylinder` – last cylinder accessed, used to simulate seek time.
- A line buffer for incoming text commands.
- A 128-byte buffer for one block of data.

Protocol and main algorithms:

For each TCP connection, the server loops: read one line, parse the command, check arguments, do the operation, send a reply, until the client closes.

Commands:

- `I`  
Client asks for geometry. Server replies with C and S on one line.
- `R c s` (read)  
Server checks that `c` and `s` are in range. If not, it returns an error.  
If valid, it computes how far the head moves (`abs(current_cylinder - c)`), sleeps for that distance times the track delay, updates `current_cylinder`, seeks to the correct offset in the backing file, reads up to 128 bytes (padding with zeros if needed), and sends a success flag plus 128 bytes of data.

- **W c s l** (write)

Server checks **c**, **s**, and **l** ( $0 \leq l \leq 128$ ). If invalid, it returns an error and ignores the write.

If valid, it reads exactly **l** bytes of data from the client into a 128-byte buffer, fills the rest with zeros, reads a trailing newline, simulates seek time as in the read case, seeks to the correct offset, writes all 128 bytes, flushes the file, and then sends success or failure.

**disk\_client** is an interactive client that sends these commands and prints the replies.

**disk\_rand** gets C and S with **I** and then runs many random **R** and **W** requests to random blocks to test the disk.

### Filesystem design (Parts 4 and 5)

The filesystem server builds a tiny filesystem on top of the disk server. It treats the disk as an array of 128-byte blocks.

Disk information:

- On startup it connects to the disk server, sends **I**, and uses C and S to compute  $\text{total\_blocks} = C \times S$ .
- Logical block numbers  $0..total\_blocks-1$  are converted back to (c, s) when reading/writing via the disk protocol.

In-memory data structures:

- Free-space map:  
**block\_used[ ]** with one entry per disk block. 0 = free, 1 = in use.
- File/directory table:  
Fixed-size array of entries. Each entry stores:
  - **used** (free/in use)
  - **is\_dir** (file or directory)
  - **parent** (index of parent directory)
  - **name** (fixed-length string)
  - **first\_block** (starting block index or -1 if no data)

- `nblocks` (number of contiguous blocks)
- `size` (file size in bytes)  
Entry 0 is the root directory “/” and has itself as parent.
- Current directory:  
`cwd_index` – index of the current directory entry.

Main algorithms:

- Format (`F`):  
Clear the entry table and free-space map, create root at entry 0, mark all blocks free, set `cwd_index` to root.
- Allocate / free blocks:  
To allocate k blocks, scan `block_used[ ]` for k consecutive zeros, mark them used, return the starting index.  
To free, given `first_block` and `nblocks`, mark that range as free.
- Create file (`C name`):  
Check that the name is not already used in the current directory, find a free entry, set it as a file with `parent = cwd_index`, no blocks, and size 0.
- Delete file (`D name`):  
Find the file in the current directory, free its blocks, and mark the entry unused.
- Write file (`W name len`):  
Find the file, read `len` bytes from the client, free any old blocks, compute how many 128-byte blocks are needed, allocate that many contiguous blocks, update `first_block`, `nblocks`, and `size`, and write the data into those blocks through the disk server.
- Read file (`R name`):  
Find the file, and if the size is > 0, read its blocks from disk and return exactly `size` bytes to the client.
- List directory (`L 0 / L 1`):  
Scan entries with `used = 1` and `parent = cwd_index`. Print their names (and, in verbose mode, type and size).

- Directories (`mkdir`, `cd`, `pwd`, `rmdir`):  
`mkdir` adds a directory entry under the current directory.  
`cd /`, `cd ..`, `cd name` update `cwd_index` to root, parent, or a child directory.  
`pwd` follows parent indices back to root and builds a path.  
`rmdir` removes a directory only if it is empty.